

Fusion redistributes per-class errors: against cRBEF it significantly gains Happy/Fear but loses Sad; against equal-weight fusion it significantly improves Sad (+3.8 pp).

Frozen M3ED formal test · mean ± SD over 3 paired seeds (cRBEF fixed checkpoint) · bootstrap on seed-43 predictions, whole-dialogue resampling (n=1,000)

