

Per class, the fusion redistributes rather than uniformly improves: significant gains over equal weight on Sad, and a trade against the fixed expert — better on Happy and Fear, worse on Sad.

Frozen M3ED formal test · mean $\pm$ SD over seeds 43/47/59 (cRBEF fixed) · 1,000 dialogue-level bootstrap resamples

