


The model degrades most when text and audio are jointly absent



Feature-level zero/mask stress; no encoder retraining; n=4201 per condition